

MidSem Project

Warp State Analysis

Team members:

Ishan Mardani (21CS01023)

Akshit Dudeja (21CS01026)

Tushar Joshi(21CS01078)

Vishnu Tirth Bysani (21CS01077)

Joshua Dias Barreto (21CS01075)

Warp States

When a kernel executes on an SM, warps of the kernel can be in different states.

We classify the warps depending on their state of execution in a given cycle:

Issued:

Warps that issue an instruction to the execution pipeline are accounted here.

It indicates the IPC of the SM and a high number of warps in this state indicate good performance

Waiting:

Warps waiting for an instruction to commit so that further dependent instructions can be issued to the pipeline are in this category.

Excess ALU (XALU):

Warps that are ready for execution of arithmetic operations, but cannot execute due to unavailability of resources are in this category. These are ready to execute warps and cannot issue because the scheduler can only issue a fixed number of instructions per cycle. Xalu indicates the excess warps ready for arithmetic execution.

Excess Memory (XMEM):

Warps that are ready to send an instruction to the Load/Store pipeline but are restricted are accounted here. These warps are restricted if the pipeline is stalled due to back pressure from memory or if the maximum number of instructions that can be issued to this pipeline have been issued. Xmem

warps represents the excess warps that will increase the pressure on the memory subsystem from the current SM.

Other:

Warps waiting on a synchronisation instruction or warps that do not have their instructions in the instruction buffer are called Others. As there is no for these warps, their requirements is unknown.

Modifications in Codebase

Issued State Counter:

```
1527     if (warp_inst_issued) {
1528         SCHED_DPRINTF(
1529             "Warp (warp_id %u, dynamic_warp_id %u) issued %u instructions\n",
1530             (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id(), issued);
1531         do_on_warp_issued(warp_id, issued, iter);
1532         // printf("@#$ISSUED\n");
1533         warp_state_counters[ISSUED]++;
1534         printf(
1535             "Warp (warp_id %u, dynamic_warp_id %u) is in ISSUED "
1536             "state\n",
1537             (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1538         customDebug.putLine("Warp (warp_id " +
1539             std::to_string((*iter)->get_warp_id()) +
1540             ", dynamic_warp_id " +
1541             std::to_string((*iter)->get_dynamic_warp_id()) +
1542             ") is in ISSUED state");
1543     }
```

Waiting State Counter:

```
1235     if (pI) {
1236         assert(valid);
1237         if (pc != pI->pc) {
1238             SCHED_DPRINTF(
1239                 "Warp (warp_id %u, dynamic_warp_id %u) control hazard "
1240                 "instruction flush\n",
1241                 (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1242             // control hazard
1243             warp(warp_id).set_next_pc(pc);
1244             warp(warp_id).ibuffer_flush();
1245             warp_state_counters[WAITING]++;
1246             printf(
1247                 "Warp (warp_id %u, dynamic_warp_id %u) is in Waiting "
1248                 "state\n",
1249                 (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1250             customDebug.putLine("Warp (warp_id " +
1251                 std::to_string((*iter)->get_warp_id()) +
1252                 ", dynamic_warp_id " +
1253                 std::to_string((*iter)->get_dynamic_warp_id()) +
1254                 ") is in Waiting state");
1255         } else {
```

```
1256         valid_inst = true;
1257         if (!m_scoreboard->checkCollision(warp_id, pI)) {
1258             SCHED_DPRINTF(
1259                 "Warp (warp_id %u, dynamic_warp_id %u) passes scoreboard\n",
1260                 (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1261             ready_inst = true;
1262
1263             const active_mask_t &active_mask =
1264                 m_shader->get_active_mask(warp_id, pI);
1265
1266             assert(warp(warp_id).inst_in_pipeline());
1267
1268             if ((pI->op == LOAD_OP) || (pI->op == STORE_OP) ||
1269                 (pI->op == MEMORY_BARRIER_OP) ||
1270                 (pI->op == TENSOR_CORE_LOAD_OP) ||
1271                 (pI->op == TENSOR_CORE_STORE_OP)) {
1272                 if (m_mem_out->has_free(m_shader->m_config->sub_core_model,
1273                     m_id)) {
```

...

```
1500     } // end of else
1501 } else {
1502     SCHED_DPRINTF(
1503         "Warp (warp_id %u, dynamic_warp_id %u) fails scoreboard\n",
1504         (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1505     // printf("!@#$WAITING\n");
1506     warp_state_counters[WAITING]++;
1507     printf(
1508         "Warp (warp_id %u, dynamic_warp_id %u) is in Waiting "
1509         "state\n",
1510         (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1511     customDebug.putLine("Warp (warp_id " +
1512         std::to_string((*iter)->get_warp_id()) +
1513         ", dynamic_warp_id " +
1514         std::to_string((*iter)->get_dynamic_warp_id()) +
1515         ") is in Waiting state");
1516 }
```

XALU State Counter:

```
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
    if (execute_on_SP) {
        m_shader->issue_warp(*m_sp_out, pI, active_mask, warp_id,
                               m_id);
        issued++;
        issued_inst = true;
        warp_inst_issued = true;
        previous_issued_inst_exec_type = exec_unit_type_t::SP;
    } else if (execute_on_INT) {
        m_shader->issue_warp(*m_int_out, pI, active_mask, warp_id,
                               m_id);
        issued++;
        issued_inst = true;
        warp_inst_issued = true;
        previous_issued_inst_exec_type = exec_unit_type_t::INT;
    } else {
        // printf("!@#$XALU\n");
        warp_state_counters[XALU]++;
        printf(
            "Warp (warp_id %u, dynamic_warp_id %u) is in X_ALU "
            "state\n",
            (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
        customDebug.putLine(
            "Warp (warp_id " +
            std::to_string((*iter)->get_warp_id()) +
            ", dynamic_warp_id " +
            std::to_string((*iter)->get_dynamic_warp_id()) +
            ") is in X_ALU state");
    }
}
```

```
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
    if (spec_pipe_avail) {
        m_shader->issue_warp(*spec_reg_set, pI, active_mask, warp_id,
                               m_id);
        issued++;
        issued_inst = true;
        warp_inst_issued = true;
        previous_issued_inst_exec_type =
            exec_unit_type_t::SPECIALIZED;
    } else {
        // printf("!@#$XALU\n");
        warp_state_counters[XALU]++;
        printf(
            "Warp (warp_id %u, dynamic_warp_id %u) is in X_ALU "
            "state\n",
            (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
        customDebug.putLine(
            "Warp (warp_id " +
            std::to_string((*iter)->get_warp_id()) +
            ", dynamic_warp_id " +
            std::to_string((*iter)->get_dynamic_warp_id()) +
            ") is in X_ALU state");
    }
}
```

```

1439         } else if ((pI->op == TENSOR_CORE_OP) &&
1440                    !(diff_exec_units && previous_issued_inst_exec_type ==
1441                      exec_unit_type_t::TENSOR)) {
1442             if (tensor_core_pipe_avail) {
1443                 m_shader->issue_warp(*m_tensor_core_out, pI, active_mask,
1444                                     warp_id, m_id);
1445                 issued++;
1446                 issued_inst = true;
1447                 warp_inst_issued = true;
1448                 previous_issued_inst_exec_type = exec_unit_type_t::TENSOR;
1449             } else {
1450                 // printf("@#XALU\n");
1451                 warp_state_counters[XALU]++;
1452                 printf(
1453                     "Warp (warp_id %u, dynamic_warp_id %u) is in X_ALU "
1454                     "state\n",
1455                     (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1456                 customDebug.putLine(
1457                     "Warp (warp_id " +
1458                     std::to_string((*iter)->get_warp_id()) +
1459                     ", dynamic_warp_id " +
1460                     std::to_string((*iter)->get_dynamic_warp_id()) +
1461                     ") is in X_ALU state");
1462             }
1463         } else if ((pI->op >= SPEC_UNIT_START_ID) &&

```

```

1387         } else if ((m_shader->m_config->gpgpu_num_dp_units > 0) &&
1388                    (pI->op == DP_OP) &&
1389                    !(diff_exec_units && previous_issued_inst_exec_type ==
1390                      exec_unit_type_t::DP)) {
1391             if (dp_pipe_avail) {
1392                 m_shader->issue_warp(*m_dp_out, pI, active_mask, warp_id,
1393                                     m_id);
1394                 issued++;
1395                 issued_inst = true;
1396                 warp_inst_issued = true;
1397                 previous_issued_inst_exec_type = exec_unit_type_t::DP;
1398             } else {
1399                 warp_state_counters[XALU]++;
1400                 printf(
1401                     "Warp (warp_id %u, dynamic_warp_id %u) is in X_ALU "
1402                     "state\n",
1403                     (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1404                 customDebug.putLine(
1405                     "Warp (warp_id " +
1406                     std::to_string((*iter)->get_warp_id()) +
1407                     ", dynamic_warp_id " +
1408                     std::to_string((*iter)->get_dynamic_warp_id()) +
1409                     ") is in X_ALU state");
1410             }
1411         } // If the DP units = 0 (like in Fermi archi), then execute DP
1412         // inst on SFU unit

```


XMEM State Counter:

```
1267
1268     if ((pI->op == LOAD_OP) || (pI->op == STORE_OP) ||
1269         (pI->op == MEMORY_BARRIER_OP) ||
1270         (pI->op == TENSOR_CORE_LOAD_OP) ||
1271         (pI->op == TENSOR_CORE_STORE_OP)) {
1272         if (m_mem_out->has_free(m_shader->m_config->sub_core_model,
1273             m_id) &&
1274             (!diff_exec_units ||
1275              previous_issued_inst_exec_type != exec_unit_type_t::MEM)) {
1276             m_shader->issue_warp(*m_mem_out, pI, active_mask, warp_id,
1277                                 m_id);
1278             issued++;
1279             issued_inst = true;
1280             warp_inst_issued = true;
1281             previous_issued_inst_exec_type = exec_unit_type_t::MEM;
1282         } else {
1283             // printf("!@$XMEM\n");
1284             warp_state_counters[XMEM]++;
1285             printf(
1286                 "Warp (warp_id %u, dynamic_warp_id %u) is in x_mem state\n",
1287                 (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1288             customDebug.putLine(
1289                 "Warp (warp_id " + std::to_string((*iter)->get_warp_id()) +
1290                 ", dynamic_warp_id " +
1291                 std::to_string((*iter)->get_dynamic_warp_id()) +
1292                 ") is in x_mem state");
1293         }
1294     } else {
1295         bool sp_pipe_avail =
1296             (m_shader->m_config->gpgpu_num_sp_units > 0) &&
1297             m_sp_out->has_free(m_shader->m_config->sub_core_model, m_id);
```


Other State Counter:

```
1192 | if (warp(warp_id).ibuffer_empty()) {
1193 |     SCHED_DPRINTF(
1194 |         "Warp (warp_id %u, dynamic_warp_id %u) fails as ibuffer_empty\n",
1195 |         (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1196 | }
1197 |
1198 | // ? why not waiting
1199 | if (warp(warp_id).waiting())
1200 |     SCHED_DPRINTF(
1201 |         "Warp (warp_id %u, dynamic_warp_id %u) fails as waiting for "
1202 |         "barrier\n",
1203 |         (*iter)->get_warp_id(), (*iter)->get_dynamic_warp_id());
1204 |
1205 | //
1206 | if (warp(warp_id).ibuffer_empty() || warp(warp_id).waiting()) {
1207 |     // printf("!@$OTHER\n");
1208 |     warp_state_counters[OTHER]++;
1209 |     customDebug.putLine("Warp (warp_id " +
1210 |         std::to_string((*iter)->get_warp_id()) +
1211 |         ", dynamic_warp_id " +
1212 |         std::to_string((*iter)->get_dynamic_warp_id()) +
1213 |         ") is in OTHER state");
1214 | }
1215 |
```

OUTPUT

We used the application Path Finder from GPU-Rodinia.

This application has three inputs: Row Length, Column Length, Pyramid Height

Input 1:

Row Length: 100

Column Length: 50

Pyramid Height: 20

Output 1:

CYCLE: 23514

ISSUED: 17632

ISSUED %: 9.529832

XALU: 2391

XALU %: 1.292300

XMEM: 447

XMEM %: 0.241597

WAITING: 101425

WAITING %: 54.818694

OTHER: 63124

OTHER %: 34.117577

Total : 185019

Input 1:

Row Length: 150

Column Length: 100

Pyramid Height: 35

Output 2:

CYCLE: 47004

ISSUED: 37338

ISSUED %: 9.982542

XALU :4950

XALU % :1.323413

XMEM :447

XMEM % :0.119508

WAITING :231269

WAITING % :61.831175

OTHER :100029

OTHER % :26.743362

Total : 374033

Input 3:

Row Length: 250

Column Length: 125

Pyramid Height: 50

Output 3:

CYCLE: 116820

ISSUED: 96410

ISSUED %: 10.338198

XALU: 12688

XALU %: 1.360554

XMEM: 894

XMEM %: 0.095865

WAITING: 612663

WAITING %: 65.696828

OTHER: 209906

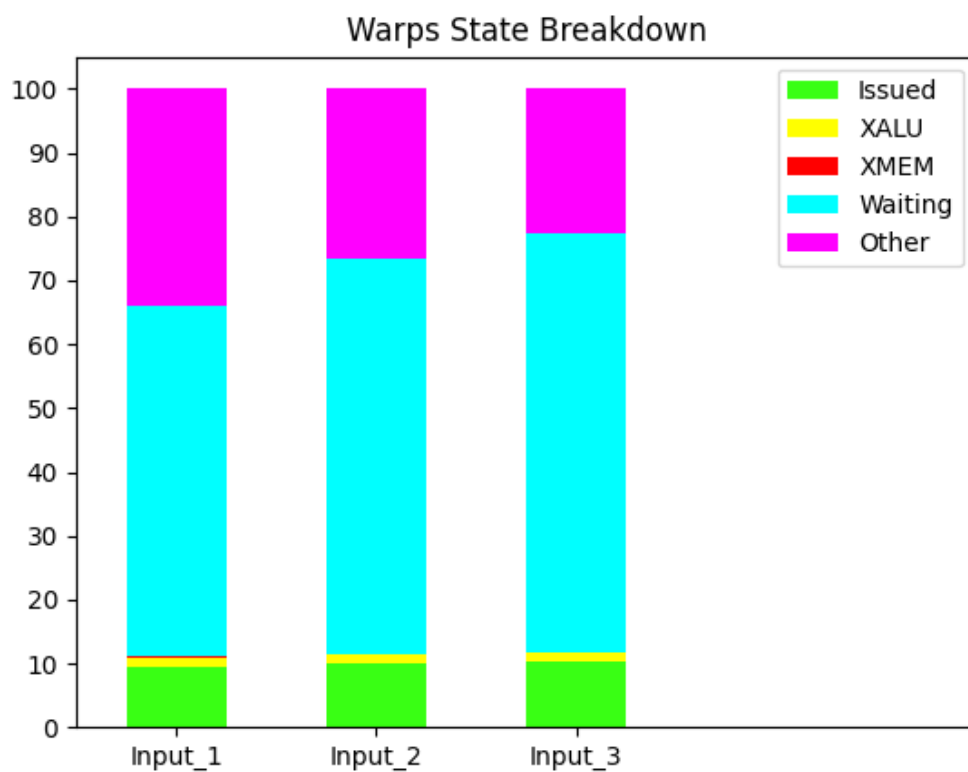
OTHER %: 22.508554

Total: 932561

Since we have 8 warps :

Cycle Count * 8 = Total State Count

Warps State Breakdown



We created a new header file warp-counters.h

```
1  #ifndef WARP_STATE_COUNTER
2  #define WARP_STATE_COUNTER
3
4  // Defining custom indexes for the counter array
5  enum counters { CYCLE, WAITING, ISSUED, XALU, XMEM, OTHER };
6
7  // Counter array size
8  #define NUM_COUNTERS (OTHER + 1)
9
10 // Declaring the counter array
11 extern unsigned long long warp_state_counters[NUM_COUNTERS];
12
13 #endif
```